

✓ Congratulations! You passed!

Grade received 100% Latest Submission Grade 100% To pass 80% or higher

[Go to next item](#)

1. At the heart of image segmentation with neural networks is an encoder/decoder architecture. What functionalities do they perform ?

1 / 1 point

- ☐ The encoder extracts features from an image and the decoder takes those extracted features, and assigns class label to the entire image.
- ☐ The decoder extracts features from an image and the encoder takes those extracted features, and assigns class labels to each pixel of the image.
- ☐ The decoder extracts features from an image and the encoder takes those extracted features, and assigns class label to the entire image.
- ☒ The encoder extracts features from an image and the decoder takes those extracted features, and assigns class labels to each pixel of the image.

✓ **Correct**
Correct!

2. Is the following statement true regarding SegNet, UNet and Fully Convolutional Neural Networks (FCNNs):

1 / 1 point

Unlike the similarity between the architecture design of SegNet & UNet, FCNNs do not have a symmetric architecture design.

- ☒ True
☐ False

✓ **Correct**
Correct!

3. What architectural difference does the *number* represent in the names of FCN-32, FCN-16, FCN-8 ?

1 / 1 point

- ☒ The *number* represents the factor by which the final pooling layer in the architecture up-samples the image to make predictions.
☐ The *number* represents the total number of convolutional layers used in the final pooling layer in the architecture to make predictions.
☐ The *number* represents the total number of filters used in the final pooling layer in the architecture to make predictions.
☐ The *number* represents the total number of pooling layers used in the architecture to help make predictions.

✓ **Correct**
Correct!

4. Take a look at the following code and select the type of scaling that will be performed

1 / 1 point

```
x = UpSampling2D(  
    size=(2, 2),  
    data_format=None,  
    interpolation='bilinear')(x)
```

- ☐ The upsampling of the image will be done by copying the value from the closest pixels.
☒ The upsampling of the image will be done by means of linear interpolation from the closest pixel values

✓ **Correct**
Correct!

5. What does the following code do?

1 / 1 point

```
Conv2DTranspose(  
    filters=32,  
    kernel_size=(3, 3)  
)
```

- ☐ It takes pixel values in the image, in a 3x3 array, and using the specified filters, creates a transpose of that array.
- ☒ It takes the pixel values and filters and tries to reverse the convolution process to return back a 3x3 array which *could* have been the original array of the image.

✓ Correct
Correct!

6. The following is the code for the *last layer* of a FCN-8 decoder. What *key change* is required if we want this to be the *last layer* of a FCN-16 decoder?

1 / 1 point

```
def fcn8_decoder(convs, n_classes):  
    ...  
  
    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(8,8),  
                                         strides=(8,8))(o)  
  
    o = (tf.keras.layers.Activation('softmax'))(o)  
  
    return o
```

- ☐ `n_classes=16`
- ☒ `kernel_size=(16, 16)`
- ☐ Using *sigmoid* instead of *softmax*.
- ☐ `strides=(16, 16)`

✓ Correct
Correct!

7. Which of the following is true about Intersection Over Union (IoU) and Dice Score, when it comes to evaluating image segmentation? (Choose all that apply.)

1 / 1 point

- ☒ For IoU the numerator is the area of overlap for both the labels, predicted and ground truth, whereas for Dice Score the numerator is 2 times that.

☒ Correct
Correct!

- ☐ For both, IoU & Dice Score the denominator is the total area of both the labels, predicted and ground truth

- ☒ Both have a range between 0 and 1

☒ Correct
Correct!

- ☐ Unlike IoU, for Dice Score the closer the value is to 0 the closer the prediction is to the ground truth.

8. Consider the following code for building the *encoder blocks* for a *U-Net*. What should this function return?

1 / 1 point

```
def unet_encoder_block(inputs, n_filters, pool_size, dropout):  
    blocks = conv2d_block(inputs, n_filters=n_filters)  
    after_pooling = tf.keras.layers.MaxPooling2D(pool_size)(blocks)  
    after_dropout = tf.keras.layers.Dropout(dropout)(after_pooling)  
  
    return # your code here
```

- ☐ *blocks*
- ☐ *after_dropout*
- ☐ *after_dropout, after_pooling* (you need to return *after_pooling* to be used in skip connections)
- ☒ *blocks, after_dropout*

☒ Correct
Correct!

9. For U-Net, on the *decoder* side you combine *skip connections* which come from the corresponding level of the *encoder*. Consider the following code and provide the missing line required to account for those skip connections with the upsampling.

1 / 1 point

(**Important Notes:** Use TensorFlow as *tf*, Keras as *keras*. And be mindful of python spacing convention, i.e (x, y) *not* (x,y))

```
def decoder_block(inputs, conv_output, n_filters, kernel_size, strides, dropout):
    upsampling_layer = tf.keras.layers.Conv2DTranspose(n_filters, kernel_size, strides = strides,
        padding = 'same')(inputs)

    skip_connection_layer = # your code here
    skip_connection_layer = tf.keras.layers.Dropout(dropout)(skip_connection_layer)
    skip_connection_layer = conv2d_block(skip_connection_layer, n_filters, kernel_size=3)

    return skip_connection_layer
```

```
tf.keras.layers.concatenate([upsampling_layer, conv_output])
```

✓ **Correct**
Correct!